

ELECTRONICS AND COMPUTER SCIENCE
Faculty of Physical and Applied Sciences
University of Southampton

Thomas A. E. Smith

taes1g09@ecs.soton.ac.uk

December 14, 2011

Intelligent Procedural Content Generation for Computer Games

Project supervisor: E. Gerding - eg@ecs.soton.ac.uk

Second examiner: C. Cirstea - cc2@ecs.soton.ac.uk

A project progress report submitted for the award of
MEng Computer Science with Artificial Intelligence

Abstract

Increasingly, as the demand for ever larger and more varied computer game environments grows, procedural content generation (PCG) is used to ensure that content remains ‘fresh’. However, many of the opportunities to use these systems to generate truly personalised content have so far been largely overlooked. When content is generated manually or algorithmically during the design phase of a game, it can only be created according to the designers’ expectations of the players’ needs. By instead generating content during the execution of the game, and using information about the player(s) as one of the system’s inputs, PCG systems should be able to produce more varied content that can be far more tailored to enhance individual players’ experiences than anything manually created. Much has been written about the generation of player models from observed data, including for the purposes of adaptivity or dynamic difficulty adjustment (DDA), and literature exists examining the problem of generating satisfying game environments via challenge adjustment. This project looks at combining these fields to create an intelligent PCG system (IPCG) that is capable of monitoring players’ progress and dynamically generating upcoming challenges to best suit their abilities.

Contents

1	Project Description	1
1.1	Requirements	1
1.1.1	Functional	1
1.1.2	Non-Functional	1
2	Project Background	3
2.1	Procedural Content Generation	3
2.2	Existing IPGC Systems	3
2.2.1	Valve’s ‘AI Director’	3
2.2.2	Bethesda’s Radiant Story	3
2.2.3	GAR’s Weapon Evolution	4
2.3	Dynamic Difficulty Adjustment	4
3	Literature Review	5
3.1	5
4	Proposed System	6
4.1	Overview	6
4.1.1	Game Base	6
4.1.2	Adaptive PCG	6
4.1.3	Player Evaluator	7
4.2	Approach	7
4.3	Justification	7
4.4	Possible Extensions	7
5	Gantt Charts	8
5.1	Completed	8
5.2	Remaining	8
	Bibliography	9

1 Project Description

The aim of this project is to investigate the use of IPGC in computer games, by looking at existing products, research in related areas and constructing a minimal prototype. Due to the increasing demand for both detail and variety within computer game environments, various aspects of in-game content are now often generated procedurally (that is, algorithmically rather than manually), using techniques that are frequently just refinements of algorithms used in the early days of computing, for games such as Elite[5]. However, one of the strengths of PCG is that (within reasonable limits) it may be performed at runtime, allowing it to also use information about the player in order to dynamically generate content on-the-fly in response to the player's actions. In general, this will involve making use of algorithms from the field of artificial intelligence in order to evaluate the information available and condense it to a form suitable for input to a PCG system; hence such systems might reasonably be termed 'intelligent' procedural content generators (IPCG). As shown later in the Literature Review and Background Research sections, IPCG techniques can be applied to many different scenarios. For the purposes of this project, the problem will be restricted to using IPCG for DDA in basic 2D platformer levels.

1.1 Requirements

The main aim of the project requirements will be to constrain the problem to an achievable scale, and inform future evaluation of the final solution. In this instance, the functional requirements are generic to any IPCG system (with refinements specific to this problem in brackets, like so), while the non-functional requirements are constraints that should serve to encourage feasibility and quality.

1.1.1 Functional

In order to properly implement IPCG, the system should:

- present the user with an interactive game environment (basic 2D platformer)
- record data on the player's behaviour (in this case, performance)
- evaluate this data according to specific criteria (a trained classifier)
- form a model of some aspect of the player (skill relative to current difficulty)
- use this model to inform further PCG activities (level chunk generation)

1.1.2 Non-Functional

In order to remain at a manageable scale, the system should:

- be written in Java
- be confined to 2D
- be presented as a basic platformer
- limit the user to move and jump actions

but also:

- remain responsive
- properly maintain the challenge of generated content

2 Project Background

2.1 Procedural Content Generation

Procedural Content Generators have been used since the early days of gaming. Well-known games such as *Elite* and *Rogue*[10] made extensive use of PCG in order to present the player with expansive game worlds far larger than could have been fully stored on the distribution media that was available at the time. As technologies improved, focus shifted more towards hand-crafted environments as it was easier to ensure that these provided value and did not feel sparse [11]. However, with the further progress of technology attention has returned to procedural generation. Modern game worlds contain vast amounts of detail, and procedural content generation algorithms are ideally suited to producing large numbers of variations on a theme, be that trees, clouds, textures, or even sounds. Producing each of these items individually by hand would take many hours of labour and much disk space, but by defining specific sub elements and assembly rules, variation can be almost endlessly reused, as in the game *Infinite Mario*[12].

2.2 Existing IPGC Systems

IPGC can be (and has been) used for a wide range of purposes, almost as varied as PCG itself. Three very different such uses are detailed below. It is unsurprising that many of the existing applications of IPGC are used to tackle some of the current key challenges in game design: maintaining players' engagement with the game via enhancing immersion and controlling flow[3].

2.2.1 Valve's 'AI Director'

One of the most well-known such applications is used in Valve's games *Left 4 Dead* and *Left 4 Dead 2*. Known as the 'AI Director', the system monitors the "emotional intensity" of each players' gameplay experience, and dynamically alters the placement of supplies and the generation of enemies of various types in order to control pacing and maintain flow. In *Left 4 Dead 2*, the Director has additional control over the structure of the level[2].

2.2.2 Bethesda's Radiant Story

Another recent example of IPGC is the Radiant Story system used in Bethesda's *Skyrim*. Rather than monitor the player's performance, it evaluates their progress

and history, and deliberately generates in-game tasks designed to force exploration of previously unknown areas; in order to increase immersion[1].

2.2.3 GAR's Weapon Evolution

Finally, the weapon evolution mechanism in the game Galactic Arms Race[6] is an unconventional application of an IPCG system, as it tracks only which weapons the players prefer, and then uses small neural networks to evolve new variations on the favourite weapons, which themselves are all simply procedurally generated particle systems[7].

2.3 Dynamic Difficulty Adjustment

Another game design concept receiving increasing attention is dynamic difficulty adjustment (DDA). Typically, challenge adjustment within video games has consisted of user choice between one or more discrete challenge settings that have been painstakingly balanced at production time. However, this solution is far from ideal - typically, if a game is begun with a certain difficulty it is difficult to later change; and this upfront decision also alienates players that are unfamiliar with the terminology or expectations, or uncertain how to classify themselves[9]. Furthermore, since game difficulty is typically a continuous function of multiple parameters, it should be possible to precisely match each player to their ideal level of challenge rather than enforcing adherence to low-resolution skill profiles. Typically, DDA is achieved by altering values that are hidden from the player, such as enemy health, accuracy, or the amount of ammo and healthkits available in the world [8]. Often, the intention is to do this invisibly, and merely ensure that the player remains optimally challenged. By manipulating values behind the scenes, it is possible to ensure that the player is neither overchallenged (leading to frustration), or underchallenged (leading to boredom)[3]. As DDA systems are given more control over additional aspects of the game environment, they can begin to enter the realm of PCG, fundamentally altering the structure and pacing of the player's experience.

3 Literature Review

Though little has been

3.1

4 Proposed System

4.1 Overview

Following the separation of concerns (SoC) design practice, the project may easily be modularised into three principle components: a ‘host’ or base module, an adaptive PCG system, and a means of taking data about the player and converting it into a model usable by the PCG. The proposed flow of information is shown in Figure 4.1.




Figure 4.1: Data flow in proposed system.

4.1.1 Game Base

A simple 2D platformer engine. This module should handle all input and rendering activity, and should follow the MVC architecture. In addition to presenting the user with the output of the IPCG system, the Base module should provide the basic input handling, physics and game functionality necessary to play the platformer, while also harvesting multiple types of information about the players performance for the Evaluator module.

4.1.2 Adaptive PCG

An adaptable 2D platforming level generator. Building on the work of Compton et al.[4], this module should maintain a context-free grammar (CFG) of obstacles available, along with weights representing the estimated challenge of each element (terminal obstacle or combination). By taking these weights into account when deriving a string of obstacles from the CFG, sequences of a desired difficulty can

be produced - or alternatively, the estimated challenge of existing sequences may be evaluated. A PCG system designed in this way should be able to generate entire levels ‘offline’, by maintaining pre-determined maxima and local variations in difficulty, but should also be capable of generating levels on-the-fly, by ensuring that short-term future difficulty levels match those requested by the Evaluator.

4.1.3 Player Evaluator

Initially, a multi-class classifier. Given the varied inputs from the Base module, the Evaluator should form a belief about the player’s skill relative to the current challenge of the level. By running the player’s data through a previously-trained classifier, this module should obtain a model that can be passed to the PCG module and acted upon.

4.2 Approach

The modules above are presented in logical order of development: none of the IPCG system will be testable without the Base program (which can be tested standalone if given a hand-crafted level), but the PCG may be run and tested using specimen models, and finally the Evaluator can be tested once the other systems are in place. The development of the classifier will involve initially collecting data on as many potentially relevant features of the player’s performance as possible, and then performing principle component analysis (PCA) upon the data-set in order to identify the maximally variant features. These can then be retained and used as the input to a K-means discretisation algorithm, which can finally be used to train a One-vs-All SVM classifier.

4.3 Justification

An alternative approach to the problem of generating a 2D platforming environment for use with DDA is presented by Sorenson et al.[13], who detail a more general top-down approach using genetic algorithms. However, the system is also more complex and provides an unneeded degree of generality for this project.

4.4 Possible Extensions

One drawback of the system as proposed is that it must still reduce all of the data about the player’s performance into a single discrete decision: whether to continue generating the level at the current difficulty, increase the difficulty, or decrease it. There is no opportunity for granularity representing player aptitude at a particular type of challenge. One possible extension would be to divide the obstacles by type (stationary hazard, timed hazard, projectile etc.), and evaluate the player’s skill on particular classes individually, then use this more detailed model to inform a slightly more sophisticated PCG module. This approach would be an ideal candidate for a collaborative filtering algorithm, which with a large enough dataset would further allow the system to predict a player’s aptitude at obstacle types that had not yet been seen by that player.

5 Gantt Charts




Figure 5.1: Completed




Figure 5.2: Remaining

Bibliography

- [1] Radiant story - the elder scrolls wiki.
- [2] M Booth. The ai systems of left 4 dead. In *Keynote, Fifth Artificial Intelligence and Interactive Digital Entertainment Conference (AIIDE '09)*, Stanford, CA, October 14 - 16, 2009. 2009.
- [3] Jenova Chen. Flow in games (and everything else). *Commun. ACM*, 50:31–34, April 2007.
- [4] K. Compton and M. Mateas. Procedural level design for platform games. 2006.
- [5] I. Bell D. Braben. Elite, 1984.
- [6] K. Stanley E. Hastings, R. Guha. Demonstrating automatic content generation in the galactic arms race video game. In *Artificial Intelligence for Interactive Digital Entertainment Conference*, 2009.
- [7] Erin J. Hastings, Ratan K. Guha, and Kenneth O. Stanley. Interactive evolution of particle systems for computer graphics and animation. *Trans. Evol. Comp*, 13:418–432, April 2009.
- [8] R. Hunicke and V. Chapman. Ai for dynamic difficulty adjustment in games. In *Papers from the 2004 AAI Workshop on Challenges in Game Artificial Intelligence*, volume Technical Report WS-04-04, Menlo Park, CA., 2004. The AAI Press.
- [9] R. Lopes and R. Bidarra. Adaptivity challenges in games and simulations: A survey. *Computational Intelligence and AI in Games, IEEE Transactions on*, 3(2):85 –99, june 2011.
- [10] K. Arnold M. Toy, G. Wichman and J. Lane. Rogue, 1983.
- [11] Michael Nitsche, Calvin Ashmore, Will Hankinson, Rob Fitzpatrick, John Kelly, and Kurt Margenau. Designing procedural game spaces: A case study. October 2006.
- [12] M Persson. Infinite mario bros! (online game), 2008.
- [13] N. Sorenson, P. Pasquier, and S. DiPaola. A generic approach to challenge modeling for the procedural creation of video game levels. *Computational Intelligence and AI in Games, IEEE Transactions on*, 3(3):229 –244, sept. 2011.